

Hand-Written notes on

Kubernetes Sidecar Containers & BusyBox

(A Simple Student Guide)

Point 1: What is a Kubernetes Sidecar Container?

1. What is a Sidecar Container?

Sometimes, along with the main container, we need an extra container to help it. This extra container is called a **Sidecar Container**. It runs alongside the main container and provides additional features to improve performance and functionality of main container without losing main containers quality and performance.

Think of it like this:

Main Container = Your application (e.g., a web server)

Sidecar Container = A helper that manages logs, monitoring, or syncing data

2. Why is a Sidecar Container needed?

If your main container is designed to do only one job but needs extra features, instead of modifying it, you can add a Sidecar Container.

Example: If a container is running a web app, but you need to store log files, you can use a separate Sidecar Container for log storage.

3. Use Cases of Sidecar Containers

- **Log Management** – To store and transfer logs
- **Security** – To enhance the security of the main container
- **Proxy Server** – To optimize networking
- **Data Backup & Sync** – To send data to cloud or storage
- **Monitoring** – To check the health of the container

4. Benefits of Sidecar Containers

- You can add new features without modifying the main container
- It improves microservices architecture
- Helps create scalable and repeatable solutions
- Makes logging, monitoring, and security easier

5. How to create and use a Sidecar Container?

You need a Pod YAML file with two containers:

1: **Main Container** (Nginx Web Server)

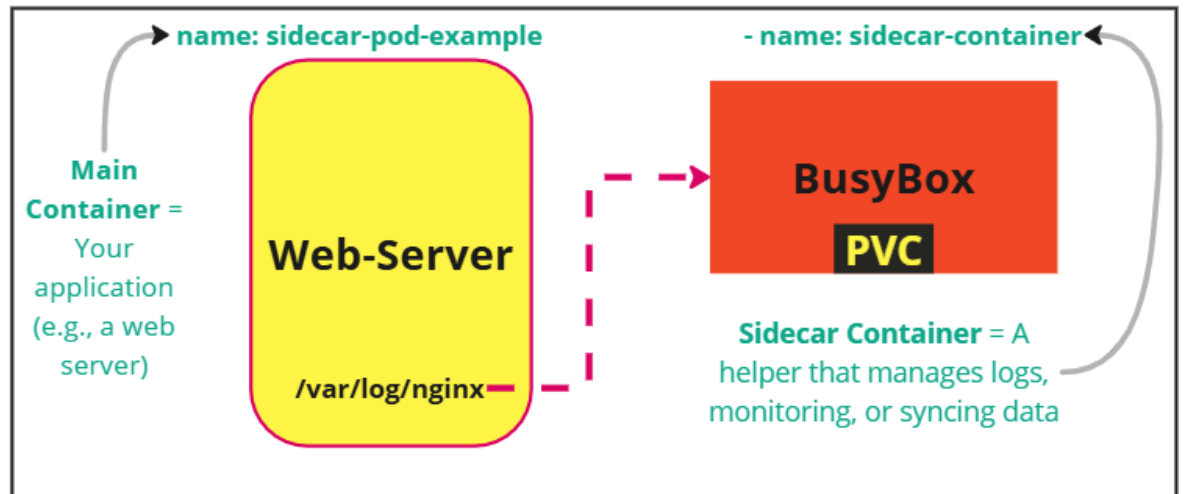
2: **Sidecar Container** (BusyBox Log Forwarder)

Sometimes, along with the main container, we need an extra container to help it. This extra container is called a **Sidecar Container**.

Definition

It runs alongside the main container and provides additional features to improve performance and functionality of main container without losing main containers quality and performance.

Diagram



Manifests

```
apiVersion: v1
kind: Pod
metadata:
  name: sidecar-pod-example
spec:
  containers:
    - name: main-container
      image: nginx
      volumeMounts:
        - name: shared-logs
          mountPath: /var/log/nginx
    - name: sidecar-container
      image: busybox
      command: ["/bin/sh", "-c", "while true; do cat /var/log/nginx/access.log; sleep 10; done"]
      volumeMounts:
        - name: shared-logs
          mountPath: /var/log/nginx
  volumes:
    - name: shared-logs
      persistentVolumeClaim:
        claimName: pvc2
```

The **main container (nginx)** writes logs to /var/log/nginx.

The **sidecar container (busybox)** reads the same logs from /var/log/nginx and prints them every 10 seconds.

pvc volume "pvc2" is using here to store data persistently

Note:

Let's understand the line below in **command** section, that runs inside the busybox container in your Kubernetes Pod:

command: ["/bin/sh", "-c", "while true; do cat /var/log/nginx/access.log; sleep 5; done"]

Breaking it Down:

1. **/bin/sh** → This starts a shell (sh) inside the container.
2. **-c** → This tells the shell to execute the following command as a script.
3. **while true; do ... done** → This is an infinite loop that keeps running continuously.

4. `cat /var/log/nginx/access.log` → Reads and prints the contents of the Nginx access log.

5. `sleep 5` → Waits for 5 seconds before running the next iteration of the loop.

What This Does in the Pod:

- The busybox container keeps reading (cat) the Nginx access logs every 5 seconds.
- This is useful for debugging because it allows you to see real-time logs.
- The loop ensures the container doesn't exit immediately after running the command once.

How does this work?

- **Main Container (Nginx)** serves the website
- **Sidecar Container (BusyBox)** continuously checks and stores Nginx logs

To apply this in Kubernetes run the manifests file:

```
kubectl apply -f sidecar.yaml
```

To check logs:

```
kubectl logs -f sidecar-example -c sidecar-container
```

Point 2: What is a BusyBox Image?

BusyBox is a very lightweight **Linux-based utility tool** that includes basic Unix commands like ls, cp, mv, echo, cat, grep, wget, etc. without any heavy program or utility tools. It is called "**The Swiss Army Knife of Embedded Linux**" because it provides multiple Unix commands in a single binary. It has 1mb size max. so this is lightweight and fast.

Use Cases of BusyBox Image

1. Testing & Debugging

You can use BusyBox for testing in Docker/Kubernetes:

```
docker run -it busybox sh
```

This opens a lightweight Linux shell where you can run Unix commands.

2. Getting Shell Access in a Container

```
kubectl run -it my-busybox --image=busybox -- sh
```

This creates a temporary pod where you can run commands inside the container.

3. Checking Network Connectivity

```
kubectl run busybox --image=busybox --restart=Never -- ping google.com
```

This checks if your Kubernetes pod is connected to the internet.

4. Data Processing & File Operations

```
kubectl run my-busybox --image=busybox --restart=Never -- cat /var/log/app.log
```

This reads log files inside the container.

5. Using BusyBox as a Sidecar Container

BusyBox can be used as a Sidecar to process data from another container:

containers:

- name: sidecar-container

image: busybox

command: ["/bin/sh", "-c", "while true; do cat /var/log/app.log; sleep 10; done"]

This continuously monitors the /var/log/app.log file.

Benefits of BusyBox Image

- **Lightweight** – Only about **1MB** in size
 - **Fast** – Loads quickly and uses fewer resources
 - **Ideal for Embedded Systems** – Works well in low-memory and low-CPU environments
 - **Complete Unix Toolset** – Includes sh, wget, ping, echo, cat, vi, etc.
-

Point 3: FAQs (Frequently Asked Questions)

1. Can we use the Main Container for Sidecar tasks?

Defiantly we can, it is possible, but it is not a good practice.

Problems

- If the **Main Container** (e.g., Nginx, PostgreSQL) handles logging, monitoring, or networking, maintenance becomes difficult.
- **The Single Responsibility Principle (SRP)** suggests that each container should do only one job.
- Updating the Main Container becomes harder if it performs multiple tasks.

2. Why is a Sidecar Container better?

- It **does not affect the Main Container**.
- The **application and Sidecar Container can be developed separately**.
- If you only need to change **log forwarding or monitoring**, you can update the Sidecar instead of modifying the Main Container.
- This approach is **better for Microservices and DevOps architectures**.

3. Why use BusyBox instead of other Base OS images?

1. Lightweight and Fast

- **BusyBox** is **less than 1MB** in size, while Base OS images (Ubuntu, Alpine, Debian) can be **100MB+**.
- A **smaller image means faster loading, less storage, and lower RAM usage**.

2. Complete Unix Command Set

- BusyBox includes essential Unix tools (ls, cat, ping, wget, grep, sh).
- We don't need a full Base OS, just basic commands.

3. Secure and Minimal

- BusyBox **only includes necessary commands**, reducing security risks.
- Other Base OS images may have unnecessary packages, increasing security vulnerabilities.

4. Best for Kubernetes and Containers

- BusyBox is a **scratch-based image** (basic and minimal), making it **perfect for containerized applications**.
- It is **ideal for Sidecar Containers, Debugging Pods, and Testing**.

4. Can we use alternatives instead of BusyBox?

Defiantly we can, but BusyBox is **the best and simplest option**. Why because of its size and minimalist nature with complete required tools options.

Alpine Linux – Lightweight (~5MB) but heavier than BusyBox

Debian Slim – Secure but larger (~20MB+)

Ubuntu Minimal – Secure and stable but heavy (~29MB+)

Scratch – **Zero-size image** (no built-in tools) but not good for debugging

If you need **Shell Access and Basic Commands in a lightweight container, BusyBox is the best choice!**

My Conclusion

- **Sidecar Containers** help improve the functionality of the Main Container without modifying it.
 - They are used for **logging, monitoring, proxy, data sync, and security**.
 - Kubernetes **Sidecar Containers** are defined in a **YAML file** with shared **volumes**.
 - **BusyBox** is a **lightweight** Linux-based image, ideal for **debugging, networking, file processing, and Sidecar Containers**.
 - It is **fast, minimal, and secure**, making it the **best choice for Kubernetes and Docker environments**.
 - Using the **Sidecar Pattern** keeps containers modular, scalable, and easy to manage.
-

Now have some practice

Practice-1: Create an Nginx web server pod with a sidecar container that prints logs every 2 seconds. This will demonstrate how a sidecar container can be used for log processing in Kubernetes.

Answer:

Step-1 Here's the YAML configuration to create a **Pod** with two containers:

1. **Nginx container** - Runs the web server and stores logs in /var/log/nginx.

2. **BusyBox sidecar container** - Reads and prints the logs from /var/log/nginx every 2 seconds.

vim sidecar-pod-example.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-sidecar-pod
spec:
  containers:
    - name: nginx-container
      image: nginx
      volumeMounts:
        - name: log-volume
          mountPath: /var/log/nginx
    - name: log-sidecar
      image: busybox
      command: ["/bin/sh", "-c"]
      args: ["while true; do cat /var/log/nginx/access.log; sleep 2; done"]
      volumeMounts:
        - name: log-volume
          mountPath: /var/log/nginx
  volumes:
    - name: log-volume
      emptyDir: {}
```

Manifests yaml file Explanation:

- **Shared Volume (emptyDir):** Both containers share /var/log/nginx using an emptyDir volume.
- **Nginx Container:** Writes access logs to /var/log/nginx/access.log.
- **BusyBox Sidecar:** Reads the logs from the shared volume and prints them every 2 seconds.

Now question may come, "If **logs are already stored in a volume**, you can view them directly from the volume mount point. So, **why do we need a sidecar container to print logs using:**

args: ["while true; do cat /var/log/nginx/access.log; sleep 2; done"]

So answer would be "You are right If you just need to view logs, use `kubectl logs` or `kubectl exec` but If you need to automatically process, filter, stream, or forward logs, a sidecar container is useful."

Step-2: See the logs from sidecar

kubectl exec -it nginx-sidecar-pod -c nginx-container -- cat /var/log/nginx/access.log

```
root@control:~# kubectl exec -it nginx-sidecar-pod -c nginx-container -- cat /var/log/nginx/access.log
127.0.0.1 - - [03/Mar/2025:07:39:33 +0000] "GET / HTTP/1.1" 200 615 "-" "curl/7.88.1" "-"
```

Real-World Projects (Sidecar Containers in Kubernetes:)

Assignment: Implementing Sidecar Containers in Kubernetes

Objective

The goal of this assignment is to understand and implement the **sidecar container pattern** in Kubernetes. Each project demonstrates how a sidecar container enhances the primary application by providing additional functionalities like log forwarding, database backup, monitoring, and dynamic configuration updates.

Project 1: Log Forwarding with Fluentd

Use Case

In microservices environments, efficient log management is crucial. This project implements a sidecar pattern where an **Nginx web server** writes logs, and a **Fluentd container** collects and forwards them to an external system.

Kubernetes Pod Manifest:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-fluentd-pod # Pod name
spec:
  volumes:
    - name: log-volume
      emptyDir: {}
  containers:
    - name: nginx # Primary container: Nginx (Web Server)
      image: nginx
      volumeMounts:
        - name: log-volume
          mountPath: /var/log/nginx
    - name: fluentd # Sidecar container: Fluentd (Log Processor)
      image: fluent/fluentd
      volumeMounts:
        - name: log-volume
          mountPath: /var/log/nginx
```

How It Works?

- Nginx writes logs to /var/log/nginx.
- Fluentd reads these logs and forwards them to an external system.

Service Manifest (ClusterIP):

```
apiVersion: v1
kind: Service
metadata:
  name: fluentd-service
spec:
  selector:
    app: nginx-fluentd
  ports:
    - protocol: TCP
      port: 24224
      targetPort: 24224
  type: ClusterIP
```

Access Fluentd Logs:

```
kubectrl port-forward svc/fluentd-service 24224:24224
```

Project 2: Automated Database Backup

Use Case

Automated backups are essential for databases running inside Kubernetes. This project introduces a **sidecar container** that periodically backs up **MySQL database data** to a shared volume.

Kubernetes Pod Manifest:

```
apiVersion: v1
kind: Pod
metadata:
  name: mysql-backup-pod
```

```

spec:
  volumes:
    - name: db-backup
      nfs:
        server: <NFS_SERVER_IP>
        path: /exported/path
  containers:
    - name: mysql
      image: mysql:5.7
      env:
        - name: MYSQL_ROOT_PASSWORD
          value: "root"
      volumeMounts:
        - name: db-backup
          mountPath: /var/backups
    - name: backup-sidecar
      image: busybox
      volumeMounts:
        - name: db-backup
          mountPath: /var/backups
      command: ["/bin/sh", "-c", "while true; do cp -r /var/lib/mysql /var/backups; sleep 3600; done"]

```

How It Works?

- MySQL stores data in its primary volume.
- The sidecar runs an automated backup process every hour.

Service Manifest (ClusterIP):

```

apiVersion: v1
kind: Service
metadata:
  name: mysql-backup-service
spec:
  selector:
    app: mysql-backup
  ports:
    - protocol: TCP
      port: 3306
      targetPort: 3306
  type: ClusterIP

```

Access MySQL Database:

```
kubectl port-forward svc/mysql-backup-service 3306:3306
```

Project 3: Application Monitoring with Prometheus

Use Case

Observability is key in modern applications. This project integrates **Prometheus as a sidecar container** to monitor application metrics in real-time.

Kubernetes Pod Manifest:

```

apiVersion: v1
kind: Pod
metadata:
  name: app-monitoring-pod
spec:
  containers:
    - name: my-app
      image: myapp:latest
      ports:
        - containerPort: 8080
    - name: prometheus-exporter
      image: prom/prometheus

```



```
ports:
  - containerPort: 9090
```

How It Works?

- The application runs on port 8080.
- The Prometheus sidecar scrapes and exposes application metrics.

Service Manifest (NodePort):

```
apiVersion: v1
kind: Service
metadata:
  name: prometheus-exporter-service
spec:
  selector:
    app: app-monitoring
  ports:
    - protocol: TCP
      port: 9090
      targetPort: 9090
      nodePort: 30090
  type: NodePort
```

Access Prometheus Dashboard:

```
kubectl get svc prometheus-exporter-service
```

Then, open:

```
http://<NODE_IP>:30090/metrics
```

Project 4: Dynamic Configuration Updates

Use Case

Many applications require **dynamic configuration updates** without restarting. This sidecar container **fetches updates periodically** and updates the config file.

Kubernetes Pod Manifest:

```
apiVersion: v1
kind: Pod
metadata:
  name: config-sync-pod
spec:
  volumes:
    - name: config-volume
      emptyDir: {}
  containers:
    - name: my-app
      image: my-app:latest
      volumeMounts:
        - name: config-volume
          mountPath: /etc/app/config
    - name: config-updater
      image: busybox
      volumeMounts:
        - name: config-volume
          mountPath: /etc/app/config
      command: ["/bin/sh", "-c", "while true; do echo 'updated' > /etc/app/config/config.yaml; sleep 60; done"]
```

How It Works?

- The application reads configuration from /etc/app/config/config.yaml.
- The sidecar updates this file every 60 seconds.

Service Manifest (ClusterIP):

```
apiVersion: v1
```

```
kind: Service
metadata:
  name: config-sync-service
spec:
  selector:
    app: config-sync
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 8080
  type: ClusterIP
```

Access Configuration Update Logs:

```
kubectl logs -f config-sync-pod -c config-updater
```

Deployment & Testing

Apply All Manifests:

```
kubectl apply -f fluentd-service.yaml
kubectl apply -f mysql-backup-service.yaml
kubectl apply -f prometheus-exporter-service.yaml
kubectl apply -f config-sync-service.yaml
```

Check Running Services:

```
kubectl get svc
```

Delete Any Service If Needed:

```
kubectl delete -f fluentd-service.yaml
kubectl delete -f mysql-backup-service.yaml
kubectl delete -f prometheus-exporter-service.yaml
kubectl delete -f config-sync-service.yaml
```

Now, verify that all your **sidecar containers** are correctly configured and functional!

Follow this channel for more: <https://www.linkedin.com/in/rakeshkumarjangid/>

