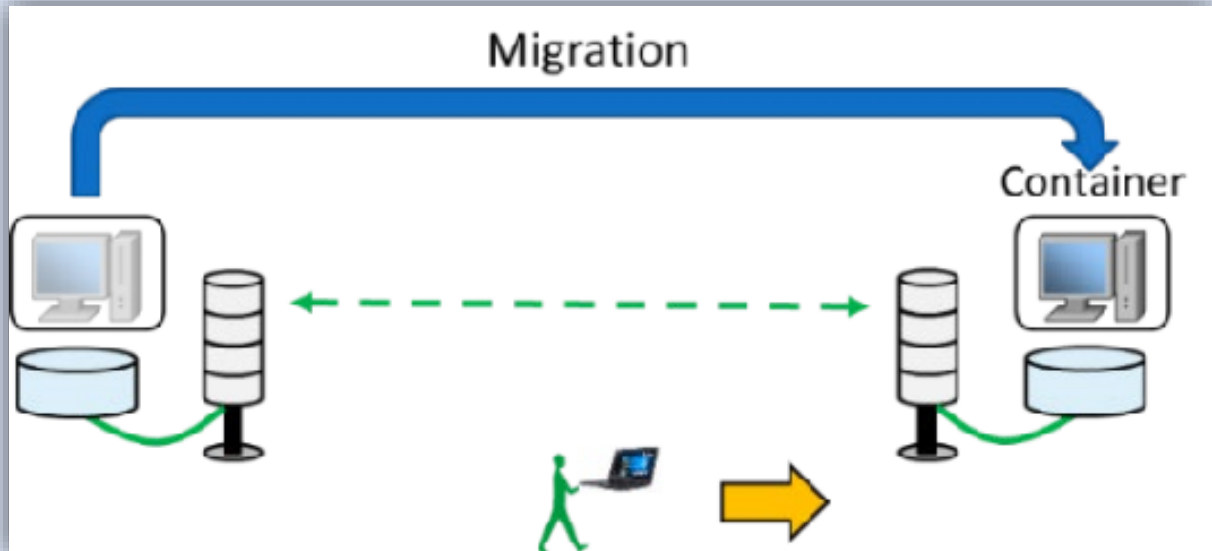


Live Migration & Checkpointing in Podman & Docker

Seamlessly Save, Transfer, and Restore Containers with Minimal Downtime



What is Checkpoint?

Checkpointing in **Podman and Docker** is a feature that allows you to **save the state of a running container** and restore it later. It uses **CRIU (Checkpoint/Restore in Userspace)** to freeze the container, store its state, and restart it when needed.

How It Works?

1. **Create a checkpoint** – The container's memory, process state, and network connections are saved.
2. **Export checkpoint** – The checkpoint files can be transferred to another system.
3. **Restore container** – The container resumes from the saved state, even on a different machine.

Use Cases

- **Live migration** – Move running containers across hosts without downtime.
- **Fault tolerance** – Restore failed containers to their last known state.
- **Reducing startup time** – Resume long-running processes instead of restarting from scratch.
- **Snapshot-based rollback** – Revert to a stable state if something goes wrong.

Advantages

- ✓ **Minimizes downtime** – Ideal for critical applications.
 - ✓ **Improves flexibility** – Containers can be moved between hosts.
 - ✓ **Enhances recovery** – Restores containers quickly after failures.
 - ✓ **Optimizes resources** – Saves CPU/memory by suspending unused containers.
-

Step 1: Install Required Packages

- **Ensure you have Podman and CRIU installed:**
 - `Yum install -y podman criu`

```
root@rhel:~# yum install -y podman criu
Updating Subscription Management repositories.
Last metadata expiration check: 0:07:51 ago on Tue 04 Mar 2025 06:26:28 AM UTC.
Package podman-4:5.2.2-13.el9_5.x86_64 is already installed.
Package criu-3.19-1.el9.x86_64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
```

- **Verify installation:**
 - `podman --version`
 - `criu --version`

```
root@rhel:~# podman --version
criu --version
podman version 5.2.2
Version: 3.19
```

Step 2: Run a Container

- **Start a container that you want to checkpoint. Example:**
 - `podman run -d --name mycontainer alpine sleep 1000`

```
root@rhel:~# podman run -d --name mycontainer alpine sleep 1000
Resolved "alpine" as an alias (/etc/containers/registries.conf.d/000-shortnames.conf)
Trying to pull docker.io/library/alpine:latest...
Getting image source signatures
Copying blob f18232174bc9 done |
Copying config aded1e1a5b done |
Writing manifest to image destination
bf116a83b38dc0a7499f1544688ae97d73fde654a3bd61c4f15cf3191a3bc3de
```

Live Migration using Checkpoint

- Check the running container:

- `podman ps`

```
root@rhel:~# podman ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED          STATUS          PORTS          NAMES
bf116a83b38d   docker.io/library/alpine:latest     sleep 1000              About a minute ago Up About a minute               mycontainer
```

```
root@rhel:~# podman top mycontainer
USER      PID      PPID      %CPU      ELAPSED      TTY      TIME      COMMAND
root      1         0          0.000      9.624714771s ?          0s        sleep 1000
```

Step 3: Checkpoint the Container

- Save the container's state using:

- `podman container checkpoint -l`

```
root@rhel:~# podman container checkpoint -l
bf116a83b38dc0a7499f1544688ae97d73fde654a3bd61c4f15cf3191a3bc3de
```

OR

- `podman container checkpoint mycontainer`

```
root@rhel:~# podman container checkpoint mycontainer
mycontainer
```

This creates checkpoint files in `/var/lib/containers/storage/overlay-containers/<container_id>/userdata/checkpoint/`.

- To verify checkpoint creation:

- `ls /var/lib/containers/storage/overlay-containers/${podman inspect --format '{{.Id}}' mycontainer}/userdata/checkpoint/`

Step 4: Export the Checkpoint (Optional - For Migration)

- To migrate a container to another system, export the checkpoint:

- `podman container checkpoint --export=/tmp/mycontainer.tar mycontainer`

```
root@rhel:~# podman container checkpoint --export=/tmp/mycontainer.tar mycontainer
mycontainer
```

Live Migration using Checkpoint

- Copy the checkpoint to another system (example using SCP):
 - `scp /tmp/mycontainer.tar user@remote_host:/tmp/`

```
root@rhel:~# scp /tmp/mycontainer.tar client1:/tmp/
mycontainer.tar
```

100% 27KB 14.8MB/s 00:00

Step 5: Restore the Container

A) Restore on the Same System

- To restore the container:
 - `podman container restore -l`
- OR
- `podman container restore mycontainer`

B) Restore on a Different System (After Importing Checkpoint)

- On the new system, import the checkpoint:
 - `podman container restore --import=/tmp/mycontainer.tar`

```
root@rhel:~# podman container checkpoint --export=/tmp/mycontainer.tar mycontainer
mycontainer
```

- Check if the container is running again:
 - `podman ps`

```
root@client1:~# podman ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
bf116a83b38d	docker.io/library/alpine:latest	sleep 1000	About a minute ago	Up 4 seconds		mycontainer

```
root@client1:~# podman top mycontainer
```

USER	PID	PPID	%CPU	ELAPSED	TTY	TIME	COMMAND
root	1	0	0.000	24.378214598s	?	0s	sleep 1000

Docker

Step 1: Run a Container

- Start a test container:
 - `docker run -d --name mygame sidhu1504/snake-game:latest`

Live Migration using Checkpoint

```
root@rhel:~# docker run --name=demo1 -d docker.io/httpd
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
Trying to pull docker.io/library/httpd:latest...
Getting image source signatures
Copying blob fdebd6c6e1b2 done |
Copying blob 7cf63256a31a done |
Copying blob d2f10b557009 done |
Copying blob 4f4fb700ef54 done |
Copying blob 38fd0d422c41 done |
Copying blob 470035b3d48f done |
Copying config 0de612e991 done |
Writing manifest to image destination
288270261ea0cd1c8ec4e4f1e46eadbf15d3d3d27e7757f7d080591815505227
```

- Check the running container:

- `docker ps`

```
root@rhel:~# docker ps
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS        NAMES
288270261ea0   docker.io/library/httpd:latest      httpd-foreground        3 seconds ago Up 3 seconds  80/tcp       demo1
```

```
root@rhel:~# docker exec -it demo1 bash
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
root@288270261ea0:/usr/local/apache2# sleep 1000
^C
root@288270261ea0:/usr/local/apache2# sleep 1000 &
[1] 89
root@288270261ea0:/usr/local/apache2# exit
exit
```

```
root@rhel:~# docker top demo1
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
USER          PID          PPID          %CPU          ELAPSED        TTY          TIME          COMMAND
root           1             0             0.000         3m55.809543704s ?            0s            httpd -DFOREGROUND
www-data       3             1             0.000         3m55.809798206s ?            0s            httpd -DFOREGROUND
www-data       4             1             0.000         3m55.809892796s ?            0s            httpd -DFOREGROUND
www-data       5             1             0.000         3m55.809954321s ?            0s            httpd -DFOREGROUND
root           89           1             0.000         31.810016934s  ?            0s            sleep 1000
```

Step 2: Checkpoint the Container

- Save the container's state:
 - `docker checkpoint create mygame checkpoint1`
- To verify checkpoint creation:
 - `ls /var/lib/docker/containers/$(docker inspect --format '{{.Id}}' mygame)/checkpoints/`

Step 3: Export the Checkpoint (Optional - For Migration)

Live Migration using Checkpoint

- To migrate the checkpoint to another system, export it:
 - `docker checkpoint create --checkpoint-dir=/tmp mygame checkpoint1`

```
root@rhel:~# docker container checkpoint demo1 -e /tmp/mycheckpoint.tar.gz
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
demo1
```

- Copy the checkpoint to another system using scp:
 - `scp -r /tmp/checkpoint1 user@remote_host:/tmp/`

```
root@rhel:~# scp /tmp/mycheckpoint.tar.gz client1:/tmp
mycheckpoint.tar.gz 100% 405KB 55.6MB/s 00:00
```

Step 6: Restore the Container

A) Restore on the Same System

- To restore the container:
 - `docker start --checkpoint checkpoint1 mygame`

B) Restore on a Different System (After Importing Checkpoint)

- On the new system, import the checkpoint and restore:
 - `docker start --checkpoint-dir=/tmp --checkpoint checkpoint1 mygame`

```
root@client1:/tmp# docker container restore -i mycheckpoint.tar.gz
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
Trying to pull docker.io/library/httpd:latest...
Getting image source signatures
Copying blob fdebd6c6e1b2 done |
Copying blob 7cf63256a31a done |
Copying blob d2f10b557009 done |
Copying blob 4f4fb700ef54 done |
Copying blob 38fd0d422c41 done |
Copying blob 470035b3d48f done |
Copying config 0de612e991 done |
Writing manifest to image destination
288270261ea0cd1c8ec4e4f1e46eadbf15d3d3d27e7757f7d080591815505227
```

- Check if the container is running:
 - `docker ps`

```
root@client1:/tmp# docker ps
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS        NAMES
288270261ea0   docker.io/library/httpd:latest      httpd-foreground        21 seconds ago Up 21 seconds 80/tcp        demo1
```

```
root@client1:/tmp# docker top demo1
Emulate Docker CLI using podman. Create /etc/containers/nodocker to quiet msg.
USER          PID          PPID          %CPU          ELAPSED        TTY          TIME          COMMAND
root          1             0             0.000         30.728564019s ?             0s            httpd -DFOREGROUND
www-data      3             1             0.000         30.728699635s ?             0s            httpd -DFOREGROUND
www-data      4             1             0.000         30.728768413s ?             0s            httpd -DFOREGROUND
www-data      5             1             0.000         30.728826337s ?             0s            httpd -DFOREGROUND
root          89            1             0.000         30.72891145s ?             0s            sleep 1000
```